

[Table of Contents](#)   Overview

- 1. The life cycle of Linux Proc...
- 2. [Zombie process](#)
  - 2.1. How zombie process o...
  - 2.2. [Recycle of zombie proc...](#)
    - 2.2.1. [The parent process ...](#)
    - 2.2.2. [The child process fi...](#)
- 3. Harm of zombie process

## Zombie Process In Linux

📅 Posted on 2016-01-16 | 📄 In [OS](#) | 👁 Visitors:

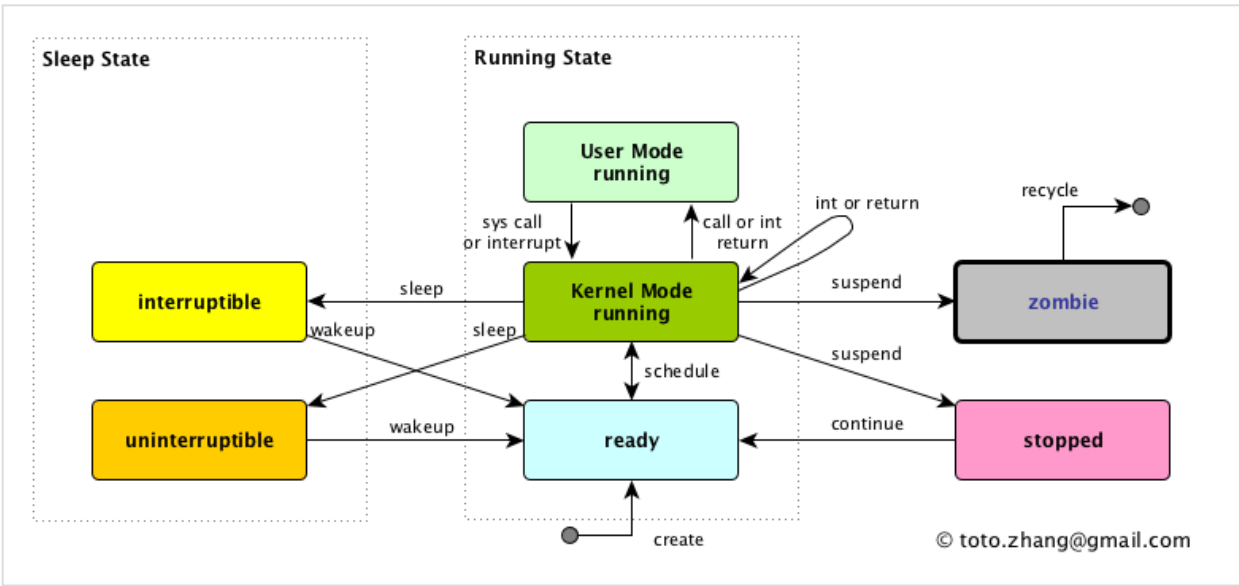
In Linux, any process that ends up running becomes a zombie process within a certain period, so a single zombie process is not inherently harmful. Only when the number of zombie processes in the system continues to accumulate and not disappear, the safety of the system will be threatened, especially in important server systems, the potential harm of the zombie process requires our special attention. So why does zombie processe exist? How does this process occur? What will happen to the system when zombie processes accumulate in large Numbers? How to avoid the potential harm of zombie process? This article discusses the above issues and briefly summarizes the process control of Linux operating system. The theory and approach to zombie processes are applicable to Solaris, BSD, and the Linux family of operating systems that conform to POSIX standards.

### The life cycle of Linux Process

In the Linux operating system, any Process is created by a previous existing process, which is called the parent process of the newly created process, and the newly created process is a child process. The only exception here is the init process, which is the first process loaded by the OS kernel. Init is the root of the process tree and the status of the init process can be seen using the pstree command.

```
1 toto@guru:~$ pstree
2 init--ModemManager--2*[{ModemManager}]
3   |
3   |--NetworkManager--dhclient
4   |
4   |-----dnsmasq
5   |
5   |-----3*[{NetworkManager}]
6   |--accounts-daemon--2*[{accounts-daemon}]
7   |--acpid
8   ...
9   |--cron
10  |--cups-browsed
11  ...
12  |--kerneloops
13  |--lightdm--Xorg--(Xorg)
14  |   |
14  |   |--lightdm--init--at-spi-bus-laun--dbus-daemon
15  |   |   |
15  |   |   |--3*[{at-spi-bus-laun}]
16  |   |   |
16  |   |   |--at-spi2-registr--(at-spi2-registr)
17  |   ...
17  |   ...
17  |   ...
18  |--wpa_supplicant
```

There are several states throughout the lifecycle of a process, including **running state**, **sleep state**, **pause state**, and **zombie state**. Among them, **running state** is subdivided into ready state, kernel running state and user running state. **sleep state** is divided into interruptible sleep state and uninterruptible sleep state. The state transitions are as follows. Any process that creates another process needs to apply to the operating system, and after the application is approved, the new created process enters the ready state. The kernel loads and runs the new process. The process switches to the kernel running state and the user running state. When the process terminates, the process enters the zombie state and stays in the zombie state until its parent process recycles it. Other states are not covered in this article.



It can be told from the state transition that the zombie state is a mandatory path that a process must go through, and the zombie process is the process in the zombie state.

### Process creation

As mentioned earlier, any process that creates another processe requires an application to the operating system through a fork system call. When a process calls fork, the operating system kernel adds a new item to its progress table and allocates resources for the new item, including memory resources, file descriptors, and so on. From the user’ s perspective, a new process is born. The new item in the process table describes all the information about the new process, and each field of the new item is described in the current Linux kernel using a struct named task\_struct, where the field pid represents the process ID and is the unique identification of the process in the kernel process table. Usually fork() is used in conjunction with the exec() family of functions. Refer man manual or [APUEv3](#) for details. The following figure describes the procedure of how a process called PIDm creates PIDx.

- Step 1. (1) process PIDm calls fork and enters the kernel for execution.
- Step 2. (2) the kernel assigns the contents of the task\_struct structure to the new process and adds this to the process table.
- Step 3. (3) the new item describes a newly assigned process PIDx.
- Step 4. (4) the fork call returns the process PIDm from the kernel with a PIDx value.
- Step 5. (4) fork the call returns the process PIDx from the kernel with a return value of 0 to distinguish between the parent process PIDm and the child process PIDx.

PIDm and PIDx hold the same memory space, file descriptor and other resources. The differences and similarities of the resources can be referred to “man fork” .

1. The life cycle of Linux Proc...
2. Zombie process
  - 2.1. How zombie process o...
  - 2.2. Recycle of zombie proc...
    - 2.2.1. The parent process ...
    - 2.2.2. The child process fi...
3. Harm of zombie process



```

1  int main(int argc, char** argv)
2  {
3      pid_t pid = Fork();
4
5      // Child, PID is 0, STEP 5
6      if (pid == 0)
7      {
8          // exec functions are called to start a executable programme.
9          printf("child process: PID = %d\n", getpid());
10         proc_child();
11         exit(0);
12     }
13
14     // Parent, PID > 0, STEP 4
15     // Parent continues
16     printf("parent process: PID = %d\n", getpid());
17     proc_parent();
18     exit(0);
19 }

```

There are many ways to end a process. For example, the process has some ways to end itself, such as returning from the main function, calling `exit()`, `_exit()`, `_Exit()` function, the last thread of the process ends or the last thread of the process calls `pthread_exit()` function, which will cause the process to exit. There are many ways for a process to be forced to end, such as receiving some signals such as `SIGKILL`, `SIGABRT`, `SIGQUIT`, `SIGINT` and so on, which lead to the passive exit of the process.

The diagram illustrates the OS kernel's process table and a process's internal state. On the left, the "OS kernel" contains a "Process Table" with columns for PID, item, PIDm, and item. The table lists PID1, PID2, ..., PIDm, and ... with corresponding items. Below this, a detailed view of a process entry shows columns for PIDx, PPIDm, details, and exit status. An arrow points from the "exit status" column to a box on the right.


On the right, the "Process PIDm" box shows the process's internal state. It contains a flowchart: "call fork" leads to "fork return child process ID PIDx", which leads to "PIDm continue", which leads to "...". To the right of this flowchart is a "Process table item" box containing "memory", "files", and "others...". Arrows point from the "Process table item" box to the "Process PIDm" box, and from the "Process table item" box to the "Process PIDm" box.

Below the "Process PIDm" box, a text box states: "The process table keeps the entries for the process PIDx, waits for the parent PIDm to read the exit state of the PIDx."

© toto.zhang@gmail.com

As can be seen from the schematic diagram of process exit, from the perspective of the kernel, the exited processes that only occupy the kernel process table item and do not occupy any system resources are zombie processes. From the user's point of view, a process that has been terminated in some way, but whose exit status has not been recycled by the parent process, is a zombie process.

The root cause of zombie process is that the parent process does not recycle the exited child process, resulting in the child process to become a zombie process. Use the following code to create a zombie process and view the status of the process through the ps command.



```
1 //parent.c slice
2 int main(int argc, char** argv)
3 {
4     int i = 0;
5     pid_t pid[MAX_CHLD_PROC_NUM];
6
7     for (i = 0; i < MAX_CHLD_PROC_NUM; i++)
8     {
9         if ((pid[i] = Fork()) == 0)
10         {
11             execve("./child", NULL, NULL);
12         }
13     }
14
15     for(;;)
16     {
17         printf("parent process: PID = %d sleep.\n", getpid());
18         sleep(10);
19         printf("parent process: PID = %d wakeup.\n", getpid());
20     }
21
22     exit(0);
23 }
```



- 1. The life cycle of Linux Proc...
- 2. **Zombie process**
  - 2.1. How zombie process o...
  - 2.2. **Recycle of zombie proc...**
    - 2.2.1. The parent process ...
    - 2.2.2. The child process fi...
- 3. Harm of zombie process

```
1 //child.c slice
2 #include <stdio.h>
3
4 int main(int argc, char** argv)
5 {
6     printf("child process: PID = %d exit.\n", getpid());
7     exit(0);
8 }
```

Compile and execute. After the parent process runs, it creates 10 child processes, and sleep forever. After the 10 child processes prints PID, the 10 child processes ends. We can see all these 10 childs become zombie state.

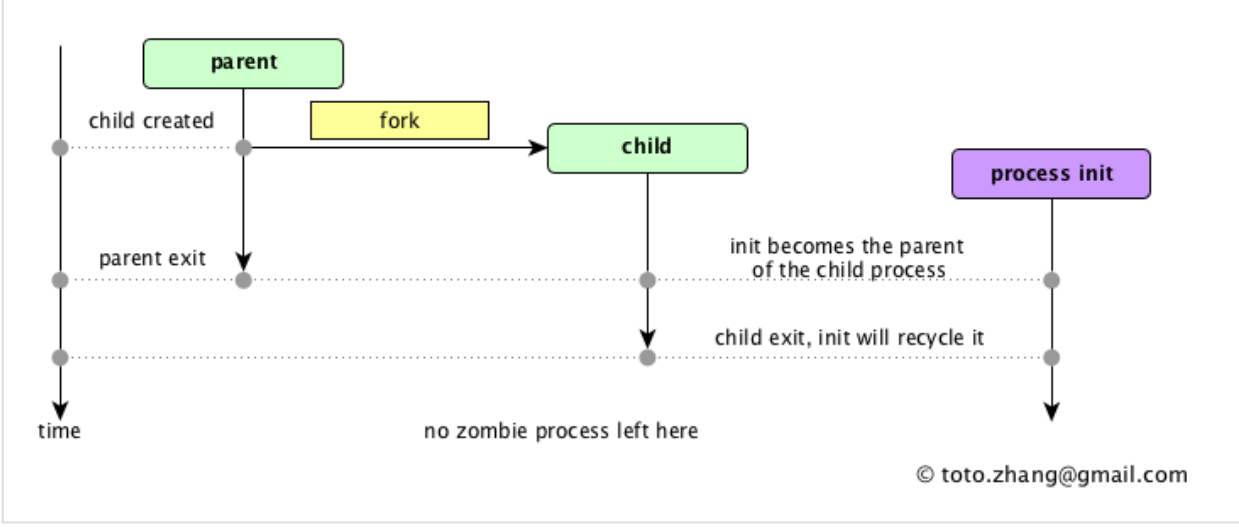
```
1 toto@guru:~$ gcc -Wimplicit-function-declaration --std=c99 parent.c -o parent
2 toto@guru:~$ gcc -Wimplicit-function-declaration --std=c99 child.c -o child
3
4 toto@guru:~$ parent
5 parent process: PID = 9286 sleep.
6 child process: PID = 9287 exit.
7 child process: PID = 9289 exit.
8 child process: PID = 9288 exit.
9 child process: PID = 9290 exit.
10 child process: PID = 9291 exit.
11 child process: PID = 9292 exit.
12 child process: PID = 9296 exit.
13 child process: PID = 9293 exit.
14 child process: PID = 9295 exit.
15 child process: PID = 9294 exit.
16 parent process: PID = 9286 wakeup.
17 parent process: PID = 9286 sleep.
18 ...
19
20 toto@guru:~$ ps -aux | grep Z
21 USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
22 toto      9287  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
23 toto      9288  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
24 toto      9289  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
25 toto      9290  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
26 toto      9291  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
27 toto      9292  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
28 toto      9293  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
29 toto      9294  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
30 toto      9295  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
31 toto      9296  0.0  0.0      0     0 pts/0    Z+   23:15   0:00 [child] <defunct
```

Recycle of zombie process

Normally, when using a multi-process model, we’ d better ensure that the parent process recycles the child processes when the child processes exit, read the exit state of child process or explicitly ignored. Be sure to avoid situations that the parent process directly ignores the child process exit status as previous example. There are two scenarios for child process recycling.

The parent process finishes running before the child process

In this case, because there is no parent process, the child process becomes the orphan process. In POSIX standard system, the process is organized as a process tree, so the orphan process will eventually become a node of the process tree, that is, a parent process must be found. At this point, the init process of the system becomes the parent of the orphan process. If the child exits at this point, the system process init recycles it. Note that the system init process does not have to be number 1 init.



Modify the child code slightly to execute the parent process again.

```
1 #include <stdio.h>
2
3 int main(int argc, char** argv)
4 {
5     while(1)
6     {
7         printf("child process: PID = %d sleep.\n", getpid());
8         sleep(10);
9         printf("child process: PID = %d wakeup.\n", getpid());
10    }
11
12    printf("child process: PID = %d exit.\n", getpid());
13    exit(0);
14 }
```

Process tree state

```
1 toto@guru:~$ pstree -p 2011
2 init(2011)─at-spi-bus-laun(2130)─dbus-daemon(2136)
3     ...
4     │   └─gnome-terminal(2661)─bash(2670)─hexo(2876)─{hexo}(2878)
5     │   ...
6     │       ...
7     │       └─parent(4259)─child(4260)
8     │                       └─child(4261)
9     │                           └─child(4262)
10    │                               └─child(4263)
11    │                                   └─child(4264)
12    │                                       └─child(4265)
13    │                                           └─child(4266)
14    │                                               └─child(4267)
15    │                                                   └─child(4268)
16    │                                                       └─child(4269)
```

Kill the parent process and view the process tree again.

```
1 toto@guru:~$ kill -9 4259
2
3 toto@guru:~$ pstree -p 2011
4 init(2011)─at-spi-bus-laun(2130)─dbus-daemon(2136)
5     │   └─{at-spi-bus-laun}(2133)
6     │   ...
7     │   └─child(4260)
8     │       └─child(4261)
9     │           └─child(4262)
10    │               └─child(4263)
11    │                   └─child(4264)
12    │                       └─child(4265)
13    │                           └─child(4266)
14    │                               └─child(4267)
15    │                                   └─child(4268)
16    │                                       └─child(4269)
```



- 1. The life cycle of Linux Proc...
- 2. Zombie process
  - 2.1. How zombie process o...
  - 2.2. Recycle of zombie proc...
  - 2.2.1. The parent process ...
  - 2.2.2. The child process fi...
- 3. Harm of zombie process

At this point, kill the child process, no more zombie process. The following is the situation of child process in the system after kill 4260 ~ 4268, I only left 4269 process stay in sleep state.

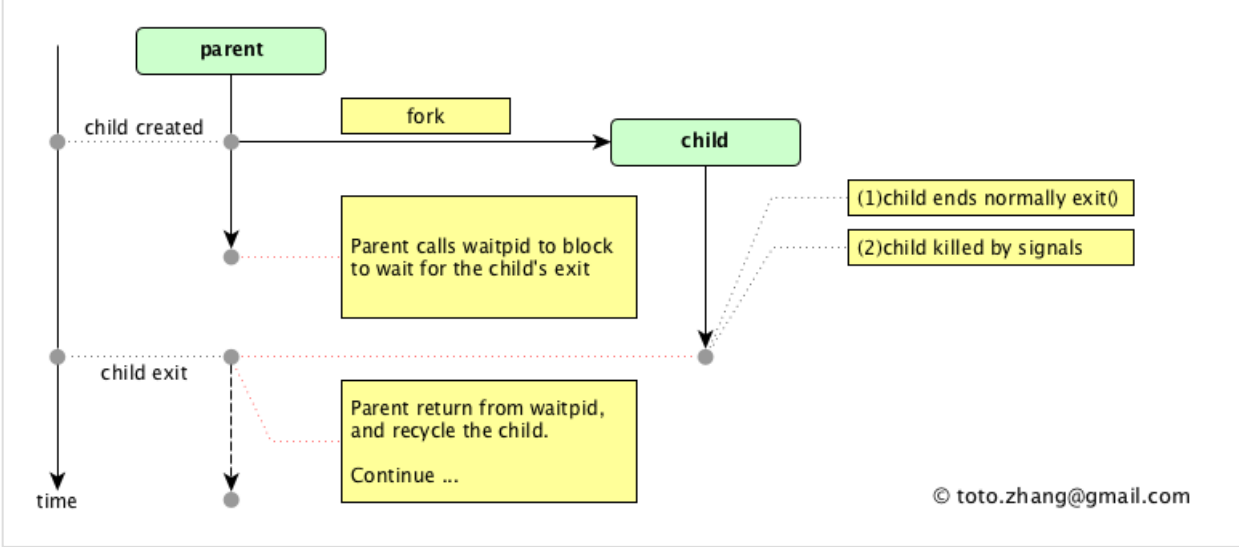
```
1 toto@guru:~$ ps -aux | grep child
2 toto      4269  0.0  0.0  4200   792 pts/9    S    07:31   0:00 [child]
```

Although init process can recycle orphaned zombie processes, when implementing multiple processes, defensive design is required to try to recycle child processes from parent processes.

The child process finishes running before the parent process

In this case, the parent process must recycle the child process itself. The system call waitpid is used for recycling. There are two methods, the first is synchronous blocking recycling, and the second is asynchronous non-blocking recycling.

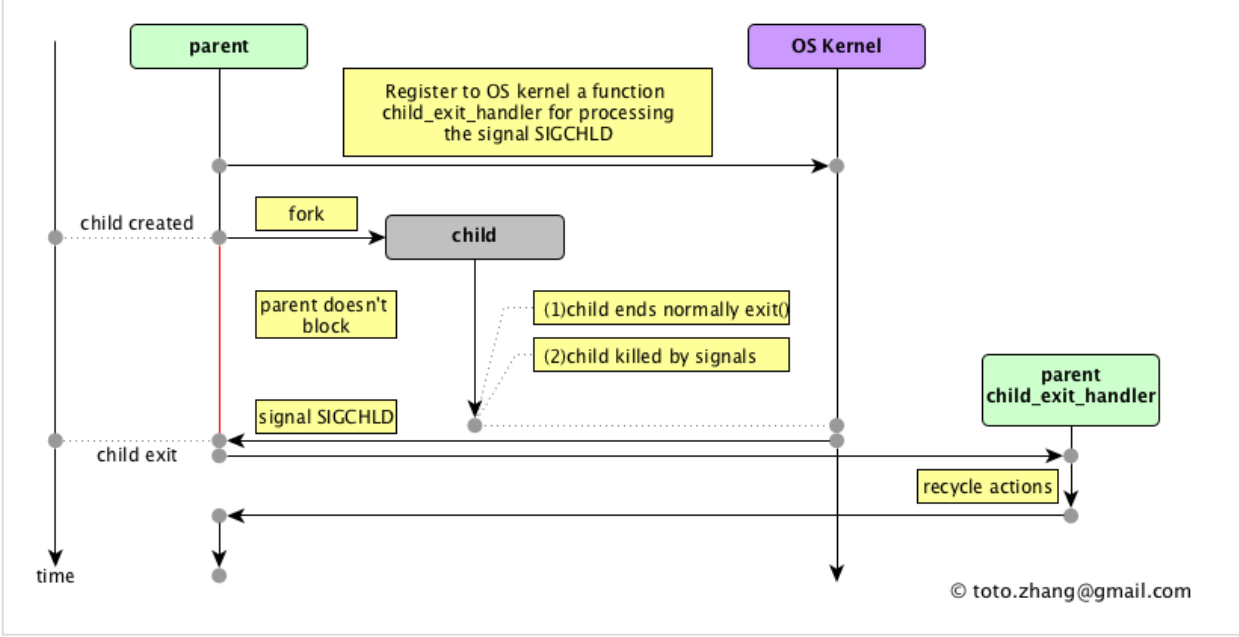
**synchronous recycling**, after the parent process creates the child process, waitpid is called and blocked to wait for all the child processes. After all the child processes finish excution, the waitpid unblocks and the parent process continues its processing until it exits.



The parent process code is modified.

```
1 //parent.c
2 int main(int argc, char** argv)
3 {
4     int status = 0;
5     pid_t ret;
6     pid_t pid[MAX_CHLD_PROC_NUM];
7
8     //10 child processes created
9     for (int i = 0; i < MAX_CHLD_PROC_NUM; i++)
10    {
11        if ((pid[i] = Fork()) == 0)
12        {
13            execve("./child", NULL, NULL);
14        }
15    }
16
17    //parent blocked to wait for all of the childs
18    while ((ret = waitpid(-1, &status, 0)) > 0)
19    {
20        if (WIFEXITED(status))
21        {
22            printf("child process %d exit with exit status %d\n", ret, WEXITSTATUS(status));
23        }
24        else if (WIFSIGNALED(status))
25        {
26            printf("child process %d killed by signal %d\n", ret, WTERMSIG(status));
27        }
28        else if (WIFSTOPPED(status))
29        {
30            printf("child process %d stoped by signal %d\n", ret, WSTOPSIG(status));
31        }
32        else
33        {
34            printf("child process %d exit unknown\n", ret);
35        }
36    }
37
38    printf("parent process exit\n");
39    exit(0);
40 }
41
```

**Asynchronous recycling**, the parent process first registers with the operating system kernel the handler when the child process exits, namely the SIGCHLD signal handler, and then continues to execute its own processing flow. Until the child process exits, the operating system kernel interrupts the parent process with signals and enters the signal processing function. After signal interrupt processing is completed, the processing flow of the parent process continues.



The parent process code is modified.

```
1 int main(int argc, char** argv)
2 {
3     pid_t pid[MAX_CHLD_PROC_NUM];
4
5     //Register OS Kernel the SIGCHLD handler
6     signal(SIGCHLD, child_exit_handler);
7
8     //10 childs
9     for (int i = 0; i < MAX_CHLD_PROC_NUM; i++)
10    {
11        if ((pid[i] = Fork()) == 0)
12        {
13            execve("./child", NULL, NULL);
14        }
15    }
16
17    //Parent continues
18    while(1)
19    {
20        printf("parent process running\n");
21        sleep(2);
22    }
23
24    exit(0);
25 }
```

Signal handler.

[Table of Contents](#) [Overview](#)

- [1. The life cycle of Linux Proc...](#)
- [2. Zombie process](#)
  - [2.1. How zombie process o...](#)
  - [2.2. Recycle of zombie proc...](#)
    - [2.2.1. The parent process ...](#)
    - [2.2.2. The child process fi...](#)
- [3. Harm of zombie process](#)

```
1 void child_exit_handler(int sig)
2 {
3     pid_t ret;
4     int status = 0;
5
6     //Parent blocked to wait
7     while ((ret = waitpid(-1, &status, 0)) > 0)
8     {
9         if (WIFEXITED(status))
10        {
11            printf("child process %d exit with exit status %d\n", ret,
12                WEXITSTATUS(status));
13        }
14        else if (WIFSIGNALED(status))
15        {
16            printf("child process %d killed by signal %d\n", ret,
17                WTERMSIG(status));
18        }
19        else if (WIFSTOPPED(status))
20        {
21            printf("child process %d stoped by signal %d\n", ret,
22                WSTOPSIG(status));
23        }
24        else
25        {
26            printf("child process %d exit unknown\n", ret);
27        }
28    }
29 }
```

### Harm of zombie process

In a server system, if a parent process continues to spawn a zombie process, it will eventually cause the kernel' s progress table to be filled up and the system will not be able to regenerate a new child process. The resulting phenomenon can be puzzling and difficult to locate. So the best way to avoid zombie processes is to make sure that when designing any multi-process system, the parent process takes the responsibilities to recycle the child processes. For the parent process that cannot be modified, in the process of operation and maintenance, some external automatic monitoring means should be used to constantly pay attention to the number of zombie processes, and restart the parent process that creates the zombie process when necessary, forcing the zombie process to be automatically recycled by the system init process.

Post author: toto

Post link: <http://totozhang.github.io/2016-01-16-linux-zombieprocess/>

Copyright Notice: All articles in this blog are licensed under [CC BY-NC-SA 4.0](#) unless stating additionally.